

Matched input



won (+)

neutral

lost (-)

same image
controlled contexts
neutral baseline

A Behavioral test

4 VLMs

Emotion
probabilities

Neutral
correction

3 readouts

B Mechanistic test

Text-trained probe
Gemma

cross-modal readout

Activation patching
Gemma · Qwen · LLaVA

image 0%
text 62-82%



Activation steering
Gemma

causal output shift